
SQL: QUERIES

Online material is available for all exercises in this chapter on the book's webpage at

<http://www.cs.wisc.edu/~dbbook>

This includes scripts to create tables for each exercise for use with Oracle, IBM DB2, Microsoft SQL Server, Microsoft Access and MySQL.

Exercise 5.1 Consider the following relations:

Student(snum: integer, sname: string, major: string, level: string, age: integer)
Class(name: string, meets_at: string, room: string, fid: integer)
Enrolled(snum: integer, cname: string)
Faculty(fid: integer, fname: string, deptid: integer)

The meaning of these relations is straightforward; for example, Enrolled has one record per student-class pair such that the student is enrolled in the class.

Write the following queries in SQL. No duplicates should be printed in any of the answers.

1. Find the names of all Juniors (level = JR) who are enrolled in a class taught by I. Teach.
2. Find the age of the oldest student who is either a History major or enrolled in a course taught by I. Teach.
3. Find the names of all classes that either meet in room R128 or have five or more students enrolled.
4. Find the names of all students who are enrolled in two classes that meet at the same time.

5. Find the names of faculty members who teach in every room in which some class is taught.
6. Find the names of faculty members for whom the combined enrollment of the courses that they teach is less than five.
7. For each level, print the level and the average age of students for that level.
8. For all levels except JR, print the level and the average age of students for that level.
9. For each faculty member that has taught classes only in room R128, print the faculty member's name and the total number of classes she or he has taught.
10. Find the names of students enrolled in the maximum number of classes.
11. Find the names of students not enrolled in any class.
12. For each age value that appears in Students, find the level value that appears most often. For example, if there are more FR level students aged 18 than SR, JR, or SO students aged 18, you should print the pair (18, FR).

Answer 5.1 The answers are given below:

1.


```
SELECT DISTINCT S.Sname
FROM   Student S, Class C, Enrolled E, Faculty F
WHERE  S.snum = E.snum AND E.cname = C.name AND C.fid = F.fid AND
       F.fname = 'I.Teach' AND S.level = 'JR'
```
2.


```
SELECT MAX(S.age)
FROM   Student S
WHERE  (S.major = 'History')
       OR S.snum IN (SELECT E.snum
                     FROM   Class C, Enrolled E, Faculty F
                     WHERE  E.cname = C.name AND C.fid = F.fid
                          AND F.fname = 'I.Teach' )
```
3.


```
SELECT  C.name
FROM    Class C
WHERE   C.room = 'R128'
       OR C.name IN (SELECT  E.cname
                     FROM    Enrolled E
                     GROUP BY E.cname
                     HAVING  COUNT (*) >= 5)
```

4.

```
SELECT DISTINCT S.sname
FROM   Student S
WHERE  S.snum IN (SELECT E1.snum
                   FROM   Enrolled E1, Enrolled E2, Class C1, Class C2
                   WHERE  E1.snum = E2.snum AND E1.cname <> E2.cname
                   AND    E1.cname = C1.name
                   AND    E2.cname = C2.name AND C1.meets_at = C2.meets_at)
```
5.

```
SELECT DISTINCT F.fname
FROM   Faculty F
WHERE  NOT EXISTS (( SELECT *
                     FROM   Class C )
                  EXCEPT
                  (SELECT C1.room
                   FROM   Class C1
                   WHERE  C1.fid = F.fid ))
```
6.

```
SELECT   DISTINCT F.fname
FROM     Faculty F
WHERE    5 > (SELECT COUNT (E.snum)
              FROM    Class C, Enrolled E
              WHERE   C.name = E.cname
              AND     C.fid = F.fid)
```
7.

```
SELECT   S.level, AVG(S.age)
FROM     Student S
GROUP BY S.level
```
8.

```
SELECT   S.level, AVG(S.age)
FROM     Student S
WHERE    S.level <> 'JR'
GROUP BY S.level
```
9.

```
SELECT   F.fname, COUNT(*) AS CourseCount
FROM     Faculty F, Class C
WHERE    F.fid = C.fid
GROUP BY F.fid, F.fname
HAVING   EVERY ( C.room = 'R128' )
```
10.

```
SELECT   DISTINCT S.sname
FROM     Student S
WHERE    S.snum IN (SELECT   E.snum
                   FROM     Enrolled E
                   GROUP BY E.snum)
```

```

HAVING COUNT (*) >= ALL (SELECT COUNT (*)
                           FROM   Enrolled E2
                           GROUP BY E2.snum ))

11.  SELECT DISTINCT S.sname
      FROM   Student S
      WHERE  S.snum NOT IN (SELECT E.snum
                           FROM   Enrolled E )

12.  SELECT  S.age, S.level
      FROM    Student S
      GROUP BY S.age, S.level,
      HAVING  S.level IN (SELECT  S1.level
                           FROM    Student S1
                           WHERE   S1.age = S.age
                           GROUP BY S1.level, S1.age
                           HAVING  COUNT (*) >= ALL (SELECT  COUNT (*)
                                                       FROM    Student S2
                                                       WHERE  s1.age = S2.age
                                                       GROUP BY S2.level, S2.age))

```

Exercise 5.2 Consider the following schema:

```

Suppliers(sid: integer, sname: string, address: string)
Parts(pid: integer, pname: string, color: string)
Catalog(sid: integer, pid: integer, cost: real)

```

The Catalog relation lists the prices charged for parts by Suppliers. Write the following queries in SQL:

1. Find the *pnames* of parts for which there is some supplier.
2. Find the *snames* of suppliers who supply every part.
3. Find the *snames* of suppliers who supply every red part.
4. Find the *pnames* of parts supplied by Acme Widget Suppliers and no one else.
5. Find the *sids* of suppliers who charge more for some part than the average cost of that part (averaged over all the suppliers who supply that part).
6. For each part, find the *sname* of the supplier who charges the most for that part.
7. Find the *sids* of suppliers who supply only red parts.
8. Find the *sids* of suppliers who supply a red part and a green part.

9. Find the *sids* of suppliers who supply a red part or a green part.
10. For every supplier that only supplies green parts, print the name of the supplier and the total number of parts that she supplies.
11. For every supplier that supplies a green part and a red part, print the name and price of the most expensive part that she supplies.

Answer 5.2 Answer omitted.

Exercise 5.3 The following relations keep track of airline flight information:

```

Flights(fno: integer, from: string, to: string, distance: integer,
        departs: time, arrives: time, price: real)
Aircraft(aid: integer, aname: string, cruisingrange: integer)
Certified(eid: integer, aid: integer)
Employees(eid: integer, ename: string, salary: integer)

```

Note that the Employees relation describes pilots and other kinds of employees as well; every pilot is certified for some aircraft, and only pilots are certified to fly. Write each of the following queries in SQL. (*Additional queries using the same schema are listed in the exercises for Chapter 4.*)

1. Find the names of aircraft such that all pilots certified to operate them have salaries more than \$80,000.
2. For each pilot who is certified for more than three aircraft, find the *eid* and the maximum *cruisingrange* of the aircraft for which she or he is certified.
3. Find the names of pilots whose *salary* is less than the price of the cheapest route from Los Angeles to Honolulu.
4. For all aircraft with *cruisingrange* over 1000 miles, find the name of the aircraft and the average salary of all pilots certified for this aircraft.
5. Find the names of pilots certified for some Boeing aircraft.
6. Find the *aids* of all aircraft that can be used on routes from Los Angeles to Chicago.
7. Identify the routes that can be piloted by every pilot who makes more than \$100,000.
8. Print the *enames* of pilots who can operate planes with *cruisingrange* greater than 3000 miles but are not certified on any Boeing aircraft.

9. A customer wants to travel from Madison to New York with no more than two changes of flight. List the choice of departure times from Madison if the customer wants to arrive in New York by 6 p.m.
10. Compute the difference between the average salary of a pilot and the average salary of all employees (including pilots).
11. Print the name and salary of every nonpilot whose salary is more than the average salary for pilots.
12. Print the names of employees who are certified only on aircrafts with cruising range longer than 1000 miles.
13. Print the names of employees who are certified only on aircrafts with cruising range longer than 1000 miles, but on at least two such aircrafts.
14. Print the names of employees who are certified only on aircrafts with cruising range longer than 1000 miles and who are certified on some Boeing aircraft.

Answer 5.3 The answers are given below:

1.


```
SELECT DISTINCT A.aname
FROM   Aircraft A
WHERE  A.Aid IN (SELECT C.aid
                FROM   Certified C, Employees E
                WHERE  C.eid = E.eid AND
                NOT EXISTS ( SELECT *
                           FROM   Employees E1
                           WHERE  E1.eid = E.eid AND E1.salary < 80000 ))
```
2.


```
SELECT   C.eid, MAX (A.cruisingrange)
FROM     Certified C, Aircraft A
WHERE    C.aid = A.aid
GROUP BY C.eid
HAVING   COUNT (*) > 3
```
3.


```
SELECT DISTINCT E.ename
FROM   Employees E
WHERE  E.salary < ( SELECT MIN (F.price)
                  FROM   Flights F
                  WHERE  F.from = 'Los Angeles' AND F.to = 'Honolulu' )
```
4. Observe that *aid* is the key for Aircraft, but the question asks for aircraft names; we deal with this complication by using an intermediate relation Temp:

- ```

SELECT Temp.name, Temp.AvgSalary
FROM (SELECT A.aid, A.aname AS name,
 AVG (E.salary) AS AvgSalary
 FROM Aircraft A, Certified C, Employees E
 WHERE A.aid = C.aid AND
 C.eid = E.eid AND A.cruisingrange > 1000
 GROUP BY A.aid, A.aname) AS Temp

```
- 5.
- ```

SELECT DISTINCT E.ename
FROM   Employees E, Certified C, Aircraft A
WHERE  E.eid = C.eid AND
        C.aid = A.aid AND
        A.aname LIKE 'Boeing%'

```
- 6.
- ```

SELECT A.aid
FROM Aircraft A
WHERE A.cruisingrange > (SELECT MIN (F.distance)
 FROM Flights F
 WHERE F.from = 'Los Angeles' AND F.to = 'Chicago')

```
- 7.
- ```

SELECT DISTINCT F.from, F.to
FROM   Flights F
WHERE  NOT EXISTS ( SELECT *
                   FROM   Employees E
                   WHERE  E.salary > 100000
                   AND
                   NOT EXISTS (SELECT *
                              FROM   Aircraft A, Certified C
                              WHERE  A.cruisingrange > F.distance
                              AND E.eid = C.eid
                              AND A.aid = C.aid) )

```
- 8.
- ```

SELECT DISTINCT E.ename
FROM Employees E
WHERE E.eid IN ((SELECT C.eid
 FROM Certified C
 WHERE EXISTS (SELECT A.aid
 FROM Aircraft A
 WHERE A.aid = C.aid
 AND A.cruisingrange > 3000)
 AND
 NOT EXISTS (SELECT A1.aid

```

```

FROM Aircraft A1
WHERE A1.aid = C.aid
AND A1.aname LIKE 'Boeing%'))

```

9.
 

```

SELECT F.departs
FROM Flights F
WHERE F.fno IN ((SELECT F0.fno
 FROM Flights F0
 WHERE F0.from = 'Madison' AND F0.to = 'New York'
 AND F0.arrives < '18:00')
 UNION
 (SELECT F0.fno
 FROM Flights F0, Flights F1
 WHERE F0.from = 'Madison' AND F0.to <> 'New York'
 AND F0.to = F1.from AND F1.to = 'New York'
 AND F1.departs > F0.arrives
 AND F1.arrives < '18:00')
 UNION
 (SELECT F0.fno
 FROM Flights F0, Flights F1, Flights F2
 WHERE F0.from = 'Madison'
 AND F0.to = F1.from
 AND F1.to = F2.from
 AND F2.to = 'New York'
 AND F0.to <> 'New York'
 AND F1.to <> 'New York'
 AND F1.departs > F0.arrives
 AND F2.departs > F1.arrives
 AND F2.arrives < '18:00'))

```
10.
 

```

SELECT Temp1.avg - Temp2.avg
FROM (SELECT AVG (E.salary) AS avg
 FROM Employees E
 WHERE E.eid IN (SELECT DISTINCT C.eid
 FROM Certified C)) AS Temp1,
 (SELECT AVG (E1.salary) AS avg
 FROM Employees E1) AS Temp2

```
11.
 

```

SELECT E.ename, E.salary
FROM Employees E
WHERE E.eid NOT IN (SELECT DISTINCT C.eid
 FROM Certified C)

```



```
AND E.salary > (SELECT AVG (E1.salary)
 FROM Employees E1
 WHERE E1.eid IN
 (SELECT DISTINCT C1.eid
 FROM Certified C1))
```

**Exercise 5.4** Consider the following relational schema. An employee can work in more than one department; the *pct\_time* field of the Works relation shows the percentage of time that a given employee works in a given department.

```
Emp(eid: integer, ename: string, age: integer, salary: real)
Works(eid: integer, did: integer, pct_time: integer)
Dept(did: integer, dname: string, budget: real, managerid: integer)
```

Write the following queries in SQL:

1. Print the names and ages of each employee who works in both the Hardware department and the Software department.
2. For each department with more than 20 full-time-equivalent employees (i.e., where the part-time and full-time employees add up to at least that many full-time employees), print the *did* together with the number of employees that work in that department.
3. Print the name of each employee whose salary exceeds the budget of all of the departments that he or she works in.

| <i>sid</i> | <i>sname</i> | <i>rating</i> | <i>age</i> |
|------------|--------------|---------------|------------|
| 18         | jones        | 3             | 30.0       |
| 41         | jonah        | 6             | 56.0       |
| 22         | ahab         | 7             | 44.0       |
| 63         | moby         | <i>null</i>   | 15.0       |

**Figure 5.1** An Instance of Sailors

4. Find the *managerids* of managers who manage only departments with budgets greater than \$1 million.
5. Find the *enames* of managers who manage the departments with the largest budgets.
6. If a manager manages more than one department, he or she *controls* the sum of all the budgets for those departments. Find the *managerids* of managers who control more than \$5 million.
7. Find the *managerids* of managers who control the largest amounts.
8. Find the *enames* of managers who manage only departments with budgets larger than \$1 million, but at least one department with budget less than \$5 million.

**Answer 5.4** Answer omitted.